

EDY NASTASE

e.nastase1@universityofgalway.ie | [LinkedIn](#) | 0894209748

Summary

On track to graduate with first-class honours in MSc Computer Science (Artificial Intelligence). Currently conducting voluntary AI Safety research through the Mentorship for Alignment Research Students (MARS) and Supervised Program for Alignment Research (SPAR) programmes. Professional experience in data analysis, machine learning, and software development.

Education

2024 - 2026 | MSc of Computer Science (AI) | University of Galway

Core Modules: Programming and Tools for Artificial Intelligence, Machine Learning, Deep Learning, Statistics for AI, Natural Language Processing, Ethics, Agents, Multi-Agent Systems and Reinforcement Learning, Knowledge Representation.

Key Projects: Scalable oversight via semi-structured reasoning (pseudo-code) to eliminate natural language ambiguities and facilitate flaw identification, ultimately enabling formal verification of LLM reasoning for safety guarantees.

2019 - 2023 | BEng in Electronic and Computing Engineering | Dublin City University

Relevant Modules: Object Oriented Programming, Computation & Simulation of Statistical Data, Semiconductor Theory and Control Systems, Numerical Problem Solving for Engineers, Operating Systems, and Algorithms for Engineers.

Research Experience

Jul 2025 - Jan 2026 | Mentorship for Alignment Researchers (MARS) | Cambridge AI Safety Hub

- Researched Scalable Oversight via AI Debate under Dr Gabriel Recchia (formerly Winton Center for Risk and Evidence Communication, University of Cambridge)
- Replicated a seminal AI Debate paper and extended it to an alternative task: designed the experimental setup, tested the new task, and integrated it within the existing codebase. The task has since been adopted in an adjacent project, demonstrating broader utility beyond the original scope.

- Paper in preparation for submission to TMLR (Mar/April 2026).
- Gained direct experience with the FindTheFlaws dataset, which establishes methodological continuity with the proposed PhD work.

Sep 2025 - Present | Supervised Program for Alignment Research (SPAR) | Kairos

- Investigating human-AI complementarity methods under Rishub Jain (Senior Research Engineer, Google DeepMind, and AlphaFold contributor), exploring how human-AI teams can outperform either alone.
- Primary developer of a sub-task delegation method: the system decomposes questions into sub-questions using LLMs, quantifies the model's confidence on each, and delegates low-confidence sub-questions to humans.
- Conducted in-depth literature review of decomposition methods, which directly informed the PhD proposal system design.
- Continuing into the next SPAR iteration (Feb-May 2026) with Dr Umang Bhatt (University of Cambridge) and Rishub Jain, building a working MVP for human-AI collaborative oversight . This component closely parallels the collaborative deliberation proposed in the PhD.

Professional Experience

August 2024 - May 2026 | Data Analyst | Vodafone Ireland

- Automation of tedious tasks such as speeding up generation of graphs, that boost my team's overall productivity.
- Delivering frequent presentations of my work to technical and management audiences, facilitating interactive discussions and consistently receiving positive feedback.
- Machine learning algorithms like linear regression on predicting numeric values or classifying data with K-nearest-neighbours.

March - August 2022 | Software Development Intern | Intel Ireland

- VPU Bootloader team: code development, testing, and technical documentation in a 20+ member team

Sept. 2020 - May 2021 | Lab Assistant Tutor | Dublin City University

- Provided help to the 1st year Engineering students during Electronics, Programming, Technical Design and Materials labs.

Awards

- **Huawei Seeds for the Future Scholarship 2022:** one of 50 selected from 500 STEM applicants across Irish universities (top 10%). Two-week intensive bootcamp covering AI, Big Data, Cloud Computing, and VR. Awarded €5,000.

- **Enterprise Ireland 24-Hour Design Challenge 2023:** 1st place out of 10 teams. Designed and pitched a noise-monitoring system for hospitality venues, judged by serial entrepreneurs and academics. Awarded €500.
- **AI Safety Fundamentals Certificate 2024:** Four-week technical alignment course from AI Safety Hungary.

Extra-Curricular Experience

2024 - Present | AI Safety Advocate | AI Safety Community

- Teaching assistant at AI Safety Global Society (<https://www.aisafety.group/>), delivering ARENA (Alignment Research Engineer Accelerator) content
- Engaged with AI Safety Dublin, by organising an online debate in September 2024 on a prominent Reinforcement Learning paper (Silver et al., Reward is Enough).
- Completed peer-reviewed AI Safety courses (i.e. AI Safety Hungary version of BlueDot's "AI Safety Fundamentals")

Sept. 2020 - May 2021 | Eng. & Comp. Faculty Representative | Dublin City University

- Actively participated in Student Council meetings by firmly expressing the concerns of the 2000 students I was representing.
- Provided detailed weekly newsletters for the students of my faculty, with the main topics of discussion revolving around opportunities and help facilities around the campus.